

Observed Recovery Paths after Depth Pruning: An OLMo Prune–Regrow Case Study

Anonymous ACL submission

Abstract

Depth-pruned language models are commonly continued-pretrained to restore aggregate language-model quality, but lower perplexity does not by itself identify which behaviors have returned. We report an observational case study of OLMo-2 prune–regrow recovery: retain a pretrained prefix, append two fresh blocks, continue pretraining, and follow likelihood and capability over the available checkpoints. The evaluation combines held-out likelihood, two MMLU interfaces, three closed-book QA tasks, and a broader likelihood suite. In the single 16-layer keep14 run, perplexity reaches 10.561 after 200k steps—still $1.428\times$ the intact model’s 7.398—while answer-letter MMLU is .319 versus .605. An intact 32-layer control remains near the base through its available 25k checkpoint, arguing against short-horizon corpus shift but not providing a 200k counterfactual. Complete-option scoring raises keep14 to .383; however, a fully random 16-layer operating point obtains nearly the same content score while remaining at chance on answer letters, exposing a fluency floor. The gap also appears on PopQA, TriviaQA, and NQ-open. A coupled construction comparison, non-contiguous ShortGPT-16, reaches .474 MMLU at 200k; it does not isolate block selection, inherited-block count, final-block retention, or fresh-tail effects. The contribution is therefore narrow: an OLMo recovery-path case study and a control bundle showing why likelihood, target capability, scoring interface, exact construction, and recovery budget should be reported separately. It does not establish seed-stable dynamics or a general law of post-pruning recovery.

1 Introduction

Removing transformer blocks and then continuing pretraining (CPT, often called “healing”) is an attractive route to smaller language models (Gromov et al., 2025; Kim et al., 2024; Men et al., 2024; Yang et al., 2024). A pruned model can improve

rapidly on held-out next-token loss, and endpoint perplexity is therefore a natural measure of whether retraining is working. It is not, however, a complete description of what the structural intervention removed or what CPT restored.

Perplexity averages local prediction loss over a corpus. Factual recall, multiple-choice answer selection, generation, and task-specific interfaces probe different observables. Compression studies have documented that aggregate language-model metrics can miss changes in parametric knowledge, downstream behavior, or safety (Namburi et al., 2023; Xu et al., 2024; Jaiswal et al., 2024). Prior depth-pruning studies also report healing curves and loss–task dissociations (Gromov et al., 2025; Kim et al., 2024; Sreenivas et al., 2024; Wibowo et al., 2025). We do not claim either phenomenon as new. Instead, we ask what can be learned from one densely measured OLMo prune–regrow path when it is read alongside several imperfect but complementary controls.

Research question. We ask: *what does continued pretraining restore in the observed OLMo prune–regrow runs, and which controls change the interpretation?* We treat recovery as a multi-axis object—held-out likelihood, target capability, scoring interface, exact construction, and training budget—rather than a scalar certified by perplexity.

We instantiate this case study on OLMo-2-1124-7B (OLMo Team et al., 2025). Prefix arms retain k pretrained decoder blocks, append two randomly initialized blocks, and continue pretraining on a fixed Dolmino corpus array and nominal schedule. The principal keep14 arm is tracked to 200k steps. Controls include the intact base, an unpruned model continued on the same corpus array and nominal schedule to an available 25k checkpoint, a frozen-prefix arm, a fully random same-shape model, and a non-contiguous ShortGPT-16 endpoint. Evaluation combines held-out likelihood, a

broad zero-shot likelihood suite, answer-letter and complete-option MMLU, and no-retrieval PopQA, TriviaQA, and NQ-open.

The central observation is bounded and recurs across the tested interfaces. In this run, keep14 reaches PPL 10.561, yet remains far below the intact model on answer-letter MMLU and all three closed-book tasks. Complete-option scoring narrows the MMLU gap, consistent with an answer-symbol or readout component, but the random-init content-score floor argues against interpreting that protocol as recovered subject knowledge. ShortGPT is substantially stronger on both PPL and MMLU, so the keep14 endpoint is not shared by every tested 16-layer construction. Coupled structural differences prevent a single-factor explanation, and all central training comparisons are single realizations.

Contributions.

- **An OLMo prune-regrow recovery path.** We jointly track in-domain likelihood, answer-letter MMLU, and broader behavior over available checkpoints, with the single-run scope explicit.
- **A complementary control bundle.** An available short-horizon full32 branch, same-shape operating points, two MMLU interfaces, closed-book generation, and a coupled ShortGPT construction comparison delimit competing interpretations without claiming clean factor isolation.
- **A case-study reporting proposal.** The measurements motivate reporting likelihood, target capability, interface, exact construction, and recovery budget as separate axes; they do not validate a universal protocol.

2 Related Work

We do not propose a pruning criterion or claim recovery-path analysis itself as new. Our contribution is an OLMo prune-regrow case study with a particular combination of controls and evaluation interfaces. Table 1 positions that combination against the nearest work.

Selecting and removing depth. ShortGPT and SLEB rank and remove redundant blocks (Men et al., 2024; Song et al., 2024); LaCo merges adjacent layers (Yang et al., 2024); BlockPruner searches at attention/MLP-block granularity (Zhong et al., 2025). Other studies show that

preferred removals depend on task and calibration choices (Siddiqui et al., 2024; Lu et al., 2024; Kim et al., 2026). These works optimize *which* structure to retain. We instead measure what different retained structures recover under CPT.

Recovery paths and retraining controls. Grovov et al. (2025) are the closest antecedent for post-healing loss-task dissociation after removing deeper layers. Shortened LLaMA reports CPT learning curves and compares CPT, LoRA, and pruned versus scratch initialization (Kim et al., 2024). Minitron likewise studies retraining trajectories, iterative pruning, task behavior, and initialization choices within a structured-pruning/distillation pipeline (Sreenivas et al., 2024). ItERABRe explicitly alternates block removal and recovery, plotting task-family trajectories and weak MMLU recovery (Wibowo et al., 2025). Thus neither trajectories nor “beyond perplexity” evaluation originate here.

Our narrower increment is the joint diagnostic package in one OLMo prune-regrow setting: an available same-corpus intact branch, same-shape operating points, answer-letter versus complete-option MMLU, closed-book generation, and a coupled construction comparison. PASER is complementary: it selects post-training data to make pruned-model recovery more efficient rather than using recovery measurements primarily for diagnosis (He et al., 2025). Sheared LLaMA, compact-model distillation, and DRPruning integrate pruning with data allocation, distillation, or robust training (Xia et al., 2024; Muralidharan et al., 2024; Deng et al., 2025). Our prefix-plus-fresh-tail recipe is not presented as competitive with these recovery algorithms.

Repair and concurrent work. LinearPatch and Prune&Comp attribute part of layer-pruning damage to cross-layer interface or magnitude mismatch and repair it with lightweight transformations or compensation (Chen et al., 2025, 2026). A decision-transition analysis studies where post-pruning performance collapse appears (Shi et al., 2026). SlimQwen and ShortOPD appeared within three months of this manuscript version and are treated as concurrent work (Tang et al., 2026; Zhang et al., 2026); they further cover matched-token initialization, progressive recovery, and recognition/generation behavior. These methods and analyses reinforce the limited novelty claim: the present paper contributes a measurement case study and control combination, not a new recovery

Study	Intervention	Adaptation	Trajectory	Loss+task	Scratch/init	Intact CPT	MMLU iface.	Closed-book	Construction
Gromov et al.	deep-block removal	CPT	partial	✓	—	—	—	—	depth/budget
Shortened LLaMA	depth pruning	CPT/LoRA	✓	✓	✓	—	—	—	ratio/method
Minitron	structured pruning	distillation/CPT	✓	✓	✓	—	—	—	iterative/shape
IterABRe	iterative block removal	iterative recovery	✓	✓	—	—	—	—	iteration
PASER	pruned-model recovery	selected CPT data	partial	✓	—	—	—	—	data policy
This study	prefix+fresh tail	CPT	✓	✓	op. point	25k only	✓	✓	coupled 16L

Table 1: **Nearest-work positioning.** Prior work already studies recovery curves, loss–task gaps, initialization, and iterative retraining. The narrower increment here is the *combination* of an OLMo prefix+fresh-tail path with an available short-horizon intact branch, same-shape operating points, two MMLU interfaces, closed-book QA, and an explicitly coupled 16-layer construction comparison. “Partial” denotes intermediate measurements that are not the sole analysis object; “op. point” is not an initialization-only ablation.

mechanism.

Evaluating compression beyond perplexity.

The Cost of Compression studies how pruning and quantization affect parametric knowledge (Namburi et al., 2023). Jaiswal et al. (2024) show that low perplexity can coexist with large deficits on knowledge-intensive compressed-LM benchmarks, and Xu et al. (2024) document divergent safety effects. These studies motivate multi-dimensional evaluation. Our additional focus is bookkeeping within one case study: within-arm checkpoints, an available intact same-corpus branch, same-shape operating points, scoring-interface comparison, closed-book generation, and an explicitly confounded construction comparison. Fragile Knowledge, Robust Instruction-Following reports selective capability effects under width pruning (Martra, 2025); it is adjacent evidence that compression can preserve some interfaces while damaging others. Multiple-choice evaluations are themselves interface-sensitive: first-token letter probabilities can disagree with generated text, leaderboard rankings change with evaluation details, and MMLU varies under answer reordering (Wang et al., 2024; Alzahrani et al., 2024; Gupta et al., 2024). Our content protocol is a diagnostic within this broader literature, not a clean isolation of knowledge.

3 Prune–Regrow–Heal Study

3.1 Intervention and model family

We study OLMo-2-1124-7B (OLMo Team et al., 2025), a 32-layer base language model without instruction tuning. For prefix arm $\text{keep} = k$, we retain the first k pretrained decoder blocks and append two randomly initialized blocks. We denote the cut location by $d_{\text{cut}} = k/32$ and resulting depth by $d_{\text{model}} = (k + 2)/32$. Embeddings, final normalization, and the output head are copied exactly. The 1B OLMo-2 model provides qualitative same-

family context, not a matched replication.

The principal arm, keep14, is therefore a 16-layer model containing 14 inherited prefix blocks and two fresh tail blocks. We additionally evaluate an influence-selected ShortGPT-16 endpoint that inherits 16 non-contiguous blocks, including the original final block. It is a construction comparison, not a selection-only ablation, because inherited count and fresh-tail construction also differ.

3.2 Healing

All inherited prefix arms continue pre-training on the DCLM portion of dolmino-mix-1124 (Soldaini et al., 2024), retokenized into 2,048-token windows. A disjoint 4096×2048 shard provides held-out likelihood. The maximum budget is 200k optimizer steps with effective batch size 128. In inherited runs, all trainable parameters use peak learning rate 2×10^{-5} ; AdamW, cosine decay, mixed precision, and reconstruction checks are detailed in Appendix B.

keep14 and ShortGPT complete 200k steps. Shallower prefix runs are evaluated at their latest retained checkpoints (keep8 121k, keep10 83.5k, and keep12 124k). These checkpoints were selected after knowledge-sensitive metrics appeared stable while PPL was still decreasing; there was no registered common stopping rule. We report literal steps and do not treat the depth ladder as a compute-matched 200k comparison.

3.3 Auxiliary operating points

Four auxiliary operating points delimit interpretation. The intact 32-layer model is continued on the same corpus array and nominal schedule through an available 25k checkpoint, measuring only short-horizon training drift without pruning. Frozen-front keep14 trains the fresh tail and non-block modules without adapting inherited blocks.

A fully random 16-layer model matches shape, data, batch, and maximum steps but uses a higher learning rate, so it is an operating point rather than an initialization-only ablation. Finally, ShortGPT provides a different construction at the same resulting depth. The study does not isolate its coupled structural differences, and current claims are limited accordingly.

3.4 Evaluation interfaces

Held-out perplexity is raw next-token loss with no chat template or inserted BOS. The principal downstream suite uses zero-shot likelihood scoring. Standard MMLU (Hendrycks et al., 2021) shows all four options and scores the single answer letter. A paired content protocol instead scores each complete option text, reporting both summed and per-token mean log-likelihood; the latter is the reported content headline. Because a random-init model obtains a substantial normalized-content score, we interpret this protocol as an exploratory interface diagnostic rather than a clean readout ablation or direct knowledge measure. Letter and content scoring change the prompt, candidate string, tokenization, and normalization together.

We additionally evaluate PopQA (Mallen et al., 2023), TriviaQA (Joshi et al., 2017), and NQ-open (Kwiatkowski et al., 2019) with a shared zero-shot, no-retrieval Question:/Answer: generation protocol. The broader likelihood suite covers HellaSwag, ARC, PIQA, WinoGrande, LAMBADA, BoolQ, OpenBookQA, CommonsenseQA, and SocialQA (Zellers et al., 2019; Clark et al., 2018; Bisk et al., 2020; Sakaguchi et al., 2021; Papperno et al., 2016; Clark et al., 2019; Mihaylov et al., 2018; Talmor et al., 2019; Sap et al., 2019). Metrics, sample counts, prompts, checkpoint provenance, and strict architecture reconstruction appear in Appendix B.

4 Recovery as a Multi-Axis Measurement

Perplexity aggregates token-level negative log-likelihood over a corpus. A knowledge-sensitive task instead asks whether a particular answer can be selected or generated under a task interface. CPT can therefore improve the former without changing the latter at the same observed rate. This does not make perplexity uninformative; it means that distributional fit is not a certificate for every behavior after a structural intervention.

This distinction motivates three study choices.

First, we retain within-arm checkpoints because a single endpoint cannot distinguish delayed from absent change over the observed budget. Second, answer-letter and complete-option scoring are reported separately because they mix competence and readout differently. Third, we compare structural constructions because the same nominal depth need not imply the same inherited computation. These are measurement requirements, not a causal theory of where capability resides.

5 Experiments

5.1 Experimental setup

All main results use OLMo-2-7B under the base-model protocol in §3. The central comparisons are the intact base, full32 continued-pretraining control, keep14+fresh2, ShortGPT-16, frozen-front, and fully random-init operating points. Detailed per-task results and all available checkpoints are in Appendix A.

5.2 Main result: the observed axes remain separated

Table 2 gives the central measurement. In the single half-depth keep14 run, the observed endpoint reaches PPL 10.561 after 200k steps—a large improvement, but still $1.428\times$ the base—and remains far below the base on both MMLU interfaces and all three closed-book tasks. The intact full32 branch remains near the base at its available step-25k checkpoint in held-out likelihood, MMLU, and closed-book evaluation. This argues against short-horizon corpus shift; it does not identify what would happen to the intact model through 200k steps.

The MMLU interface matters but does not erase the result. Scoring complete option text raises keep14’s above-chance recovery relative to answer-letter scoring, which is consistent with a content-to-symbol or readout contribution. Yet fully random initialization attains nearly the same normalized-content score as frozen-front despite chance-level answer letters. The content protocol therefore contains a large fluency floor; neither interface alone is identified with stored knowledge.

The separation is not specific to multiple choice. At 200k, keep14 reaches only .142 PopQA answer containment, .294 TriviaQA exact match, and .060 NQ-open exact match, compared with .257, .636, and .205 for the intact base. These factual recall tasks use a shared no-retrieval generation inter-

Operating point	PPL↓	MMLU letter	MMLU content	PopQA	TriviaQA	NQ-open
Full base (32L)	7.398	.605	.471	.257	.636	.205
Full32 continued@25k [†]	7.670	.588	.466	.228	.572	.158
keep14 + fresh2, 200k	10.561	.319	.383	.142	.294	.060
ShortGPT-16, 200k	9.780	.474	.401	—	—	—
Frozen inherited front, 200k	12.797	.262	.360	.128	.248	.050
Fully random init, 200k	11.498	.247	.360	.111	.209	.063

Table 2: **Likelihood and knowledge-sensitive evaluation recover differently in the observed OLMo operating points.** MMLU letter scores the answer symbol; MMLU content scores the complete option text using mean log-likelihood per continuation token, the reported content headline. PopQA reports normalized-answer containment; TriviaQA and NQ-open report exact match. Full32 continued pretraining stays near the base through 25k, whereas every tested half-depth point remains substantially below it. The random-init content score exposes a fluency floor, so content accuracy is not interpreted as recovered knowledge. Rows are single training runs; missing ShortGPT closed-book cells were not evaluated. [†] All full32 cells use the same intact 32-layer step-25k checkpoint, which is a short-horizon rather than 200k control. Letter values here are rounded headline aggregates; the common interface rerun in Table 16 differs slightly.




Figure 1: **Late keep14 healing continues on both axes, but the remaining capability gap is large.** From 128k to 200k, PPL improves from 10.826 to 10.561 and MMLU from .3012 to .3191; the intact base remains at .6053 MMLU. The figure supports slow late recovery, not a claim that MMLU is numerically constant.




Figure 2: **Available operating points are not summarized by PPL alone.** keep14, ShortGPT, frozen-front, and random-init are single-run 200k points; shallower prefix rows use unequal latest retained checkpoints. Filled red points are controls or ShortGPT; the dashed line connects the prefix ladder only as a visual guide. The plot is neither a matched-PPL nor a matched-compute analysis.

face, while full32 continued pretraining preserves much more of the base performance. Together with the two MMLU protocols, they argue against an answer-letter-only explanation.

Figure 2 makes the measurement issue visible. PPL and MMLU broadly co-vary, and no compressed arm reaches base PPL. Within these heterogeneous operating points, however, PPL alone does not recover the full ordering: random-init has lower PPL than frozen-front but lower MMLU. Because learning rate, trainable modules, structure, and checkpoint budget vary, this is a descriptive counterexample to a one-number summary, not a matched-quality estimate or evidence of global uncorrelatedness.

5.3 Controlled comparisons

Same-shape initialization and adaptation. At 200k, keep14, frozen-front, and fully random init share the 16-layer shape. Their metric rankings disagree: frozen-front has worse PPL than random init but higher answer-letter MMLU, while train-all keep14 is strongest on both. Paired item-level tests on a common rerun quantify the keep14 advantage over each same-shape operating point, but learning rate and trainable-module differences prevent an initialization-only or adaptation-only causal claim. Appendix Table 15 reports exact McNemar tests and paired-bootstrap intervals.

Full-model corpus control. Through the available 25k checkpoint, continuing the intact 32-layer model leaves PPL and both MMLU interfaces close to the original base and preserves most PopQA, TriviaQA, and NQ-open performance. Thus short-horizon corpus shift is insufficient to explain the large half-depth gap. This control is not step- or token-exposure-matched to keep14 at 200k, so we do not treat it as an exact long-budget counterfactual.

Scoring-interface diagnostic. Across nine evaluated checkpoints or operating points, complete-option MMLU is less depth sensitive than answer-letter MMLU. ShortGPT remains stronger than the prefix arms under both interfaces, showing that the observed endpoint differs across constructions. The comparison changes prompt and candidate scoring together and has no paired uncertainty analysis, so it is exploratory rather than a clean interface ablation. The full table and protocol definition are in Appendix Table 16.

6 Analysis

6.1 Recovery varies across knowledge domains

Figure 3 separates absolute accuracy from recovery relative to each group’s base margin. keep14 is strongest on the heterogeneous Other group (.377; 29.1% chance-adjusted recovery) and weakest on STEM (.294; 15.6%) and Humanities (.299; 16.2%). Social Science has higher raw accuracy (.335) but only 18.6% recovery because its base is much stronger. ShortGPT improves every group, most notably Other (.574) and Social Science (.554), yet still remains below full32.

Subject-level examples show that the pattern is not a single-domain artifact. Among high-sample subjects, keep14 retains relatively more on miscellaneous ($n=783$: .457 versus .803 base) and marketing ($n=234$: .440 versus .825), but little on professional law ($n=1,534$: .272 versus .460) and high-school statistics ($n=216$: .213 versus .514). These examples are descriptive: group membership mixes base difficulty, answer format, corpus exposure, factual recall, and reasoning. They do not establish that a particular knowledge type is stored in the removed blocks.

6.2 Observed trajectories differ across two prefix arms

keep8 and keep14 provide complementary within-arm trajectories. keep8’s PPL falls smoothly while aggregate MMLU changes from .2463 at 44k to .2535 at 121k; keep14 gains 1.79 points from .3012 at 128k to .3191 at 200k while PPL improves by 0.265 (Figure 1). A common per-item rerun estimates the keep14 change as 1.68 points (95% CI [1.08, 2.29]). These intervals are conditional on the realized checkpoints, and the two arms do not share a common budget, starting quality, or seed design. The supported claim is only that the two observed trajectories differ; they do not identify an inherited-depth effect or eventual convergence.

6.3 A coupled 16-layer construction comparison

ShortGPT and keep14 end at 16 blocks and complete the same 200k target budget, but differ along four coupled dimensions: ShortGPT inherits 16 rather than 14 pretrained blocks, selects them non-contiguously, retains original block 31, and uses no fresh tail. keep14 retains a contiguous prefix and appends two randomly initialized blocks. Embeddings, final normalization, and output heads are copied from the same base and therefore do not explain the difference.

The resulting single-run endpoint gap is large: ShortGPT reaches .474 MMLU versus .319 for keep14 while also obtaining lower PPL. Thus these two constructions show that nominal depth alone does not fix a unique observed endpoint. The comparison does *not* show that block 31 is uniquely responsible, that non-contiguous selection is optimal, or that fresh tails are harmful. Additional matched structures would be needed to test these hypotheses separately.

6.4 Scope checks, not replications

A more compressed OLMo-2-1B arm exhibits the same qualitative pattern of falling PPL with near-chance MMLU, providing same-family context rather than an independent replication. An existing Qwen3-8B endpoint supplies a limited cross-family check: a 14/36-layer prefix+fresh-tail model remains at .2495 MMLU after 200k steps, compared with .7297 for its base (Appendix Table 9). Qwen retains a smaller depth fraction than the OLMo keep14 recipe (39% versus 44%), which is one plausible contributor to its larger deficit, but the




Figure 3: **MMLU recovery is heterogeneous across broad domains.** Accuracy is sample weighted over the standard four MMLU groups. Recovery is chance adjusted, $(a - 0.25)/(a_{\text{base}} - 0.25)$. Full32 remains close to the base in every group; ShortGPT consistently exceeds keep14, but neither half-depth arm closes the intact-model gap.




Figure 4: **The shallow keep8 arm improves in likelihood and ordinary tasks through 44k.** The retained 44k and 121k aggregate checkpoints both remain near MMLU chance; the comparison is descriptive because the available run lacks seed replication.

corpus, architecture, and missing content/closed-book controls prevent attribution. We therefore treat Qwen as directional evidence, not a matched replication.

Arm	Step	Inh./fresh	PPL↓	core6	MMLU
full base (32L)	—	32/0	7.398	.7037	.6053
keep8+fresh2 (10L)	121k	8/2	13.333	.5238	.2535
keep10+fresh2 (12L)	83.5k	10/2	12.816	.5303	.2718
keep12+fresh2 (14L)	124k	12/2	11.443	.5669	.2752
keep14+fresh2 (16L)	200k	14/2	10.561	.5938	.3191
ShortGPT-16	200k	16/0	9.780	.6215	.4739

Table 3: **Available depth and construction endpoints.** core6 averages six commonsense/reasoning tasks; Appendix Table 10 gives all cells. ShortGPT retains layers [0–12, 16, 17, 31] and matches keep14 in total depth and completed 200k training, but differs in inherited count and fresh-tail construction. The prefix ladder is not compute matched: shallower runs use unequal latest retained checkpoints, whose PPL was still decreasing. Every row is one realized training run or base checkpoint.

7 Discussion

What does healing restore? In these observed runs, the combined evidence argues against treating recovery as a single scalar. PPL measures how well the healed model fits held-out next-token events; it improves substantially in every inherited-prefix arm. That improvement is useful—the models recover fluent continuation and several commonsense tasks—but it does not summarize answer selection or closed-book recall by itself. The endpoint scatter makes the distinction concrete: models with similar PPL can occupy meaningfully different MMLU regimes, and random initialization can outrank a frozen inherited front in PPL while underperforming it on answer letters. We therefore interpret healing as partial recovery of distributional modeling plus a task-dependent remainder in this case study, not as a binary return of the original model or a matched-PPL law.

Interface sensitivity is substantial but insufficient. Complete-option MMLU recovers more smoothly than answer-letter MMLU, showing substantial protocol sensitivity that is consistent with an answer-symbol component. The protocols also change the prompt, candidate continuation, tokenization, and normalization, so they do not isolate mapping loss. The random-init fluency floor and the independent PopQA, TriviaQA, and NQ-open gaps argue against a formatting-only account. This distinction matters for evaluation: a single multiple-choice protocol can overstate damage when its interface is brittle, whereas a content-likelihood protocol can overstate recovery when generic fluency makes plausible options easy to score. Agreement across interfaces is more informative than designating either as ground truth.

Different constructions yield different endpoints. At the same nominal 200k step budget, the keep14–ShortGPT comparison shows that these two half-depth constructions are not one operating point. ShortGPT yields a much stronger likelihood–capability endpoint than the tested prefix-plus-fresh-tail recipe. The result does not identify whether the advantage comes from the original final block, two additional inherited blocks, non-contiguous selection, or avoiding fresh layers. It shows that conclusions drawn from one prefix construction should not be generalized to every 16-layer construction; isolating the responsible factor requires additional matched structures.

Which behaviors remain difficult? The residual deficit spans factual recall (PopQA, TriviaQA, and NQ-open), broad academic knowledge (all four MMLU groups), and some answer-interface behavior. In contrast, several continuation-style common-sense tasks recover more readily. This taxonomy is behavioral rather than anatomical. Knowledge-neuron and causal tracing studies motivate asking whether particular feed-forward computations support factual recall (Dai et al., 2022; Meng et al., 2022), but neither our subject groups nor ShortGPT’s retention of block 31 identifies the missing substrate. Structure-isolation and causal interventions are required before making that claim.

A reporting proposal from this case study. The measurements motivate reporting likelihood, target capability, interface, recovery compute, and exact structural construction together. An intact model continued on the same corpus can bound drift only

over its observed horizon; same-shape operating points can expose alternative explanations only to the extent that learning rate and trainable modules are matched.

For a more interpretable recovery claim: (1) evaluate an independent target capability; (2) test a second scoring or generation interface; (3) state the horizon and matching of any continued-full-model control; and (4) report the exact structure and healing budget.

This proposal does not replace efficiency or endpoint-quality comparisons and is not validated across methods or model families. Its purpose is to prevent a low PPL alone from being interpreted as evidence that a target behavior has returned.

8 Conclusion

This OLMo case study records a recovery path rather than proposing a pruning method. In the single principal prefix run, in-domain likelihood improves substantially while answer-letter MMLU and closed-book QA remain far below the intact model. An available short-horizon full32 branch, an exploratory MMLU interface comparison, same-shape operating points, and a coupled ShortGPT construction comparison narrow the interpretation without isolating corpus, initialization, adaptation, or structural causes. Perplexity remains essential for distributional fit, but within these observed interventions it is not a sufficient one-number summary of target-behavior recovery. The practical proposal is to report likelihood, capability, interface, exact construction, and recovery budget separately, with seed and matching limitations explicit.

Limitations

This is an observational case study of one OLMo-2-7B prune–regrow recipe. keep14, ShortGPT, and the same-shape points are single runs; item-level intervals do not measure variation from initialization, data order, optimization, or block selection. Shallower arms stop at unequal retained checkpoints while PPL is still falling, and full32 stops at 25k. There was no registered common stopping rule, so neither the depth ladder nor full32 is a step-, token-, or FLOP-matched 200k counterfactual.

Random-init uses a higher learning rate, frozen-front changes the trainable set, and ShortGPT changes inherited count, selected layers, final-block retention, and fresh-tail use. These are operating points, not factor-isolating ablations. The two MMLU protocols also change prompt, candidate, tokenization, and normalization; closed-book

scores add generation and normalization choices. Only selected answer-letter contrasts have paired intervals.

The study covers English base models, one mixture, and mainly 7B scale. PPL is in-domain; no contamination audit or out-of-domain PPL is reported. We also do not report latency, throughput, memory, or recovery FLOPs. Exact reproduction is limited by unset training seeds, an unrecorded resumed data-loader offset, incomplete compute accounting, and no frozen runnable artifact. Layer-wise readouts are relegated to Appendix C and carry no mechanistic claim.

Ethical Considerations

This work studies how model compression via layer pruning and continued pretraining affects evaluated capability; it does not add a safety mechanism and changes none of the underlying model’s risks such as hallucination, biased generation, or unsafe completion. The observed PPL–MMLU dissociation carries a practical caution: a compressed model that looks healthy on aggregate language-modelling metrics may still perform poorly on a target benchmark, and deployments should evaluate the capabilities they intend to preserve rather than rely on perplexity alone.

We collect no new human-subject data and recruit no annotators. All experiments use previously released artifacts—the OLMo-2 base model, the Dolmino corpus, and standard public benchmarks—under their respective licenses. Training and large-scale evaluation consume energy. We reuse a common training recipe and report checkpoint step explicitly rather than presenting unequal-step results as compute matched. Any release should preserve the licenses and usage restrictions of the underlying model, corpus, and benchmarks. An anonymous artifact should include training/evaluation code, exact configs, environment locks, checkpoint hashes, per-item prediction files where redistribution permits, and the actual data-order/resume behavior; upstream model weights and benchmark text should remain with their original providers.

References

Norah Alzahrani, Hisham Abdullah Alyahya, Yazeed Alnumay, Sultan Alrashed, Shaykhah Alsubaie, Yusef Almushaykeh, Faisal Mirza, Nouf Alotaibi, Nora Altwairesh, Areeb Alowisheq, M. Saiful Bari, and Haidar Khan. 2024. [When benchmarks are tar-](#)

[gets: Revealing the sensitivity of large language model leaderboards](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*.

Nora Belrose, Igor Ostrovsky, Lev McKinney, Zach Furman, Logan Smith, Danny Halawi, Stella Biderman, and Jacob Steinhardt. 2023. [Eliciting latent predictions from transformers with the tuned lens](#). *arXiv preprint arXiv:2303.08112*.

Yonatan Bisk, Rowan Zellers, Ronan Le Bras, Jianfeng Gao, and Yejin Choi. 2020. PIQA: Reasoning about physical commonsense in natural language. In *AAAI Conference on Artificial Intelligence*.

Xinrui Chen, Haoli Bai, Tao Yuan, Ruikang Liu, Kang Zhao, Xianzhi Yu, Lu Hou, Tian Guan, Yonghong He, and Chun Yuan. 2025. [A simple linear patch revives layer-pruned large language models](#). In *Advances in Neural Information Processing Systems*.

Xinrui Chen, Hongxing Zhang, Fanyi Zeng, Yongxian Wei, Yizhi Wang, Xitong Ling, Guanghao Li, and Chun Yuan. 2026. [Prune&comp: Free lunch for layer-pruned LLMs via iterative pruning with magnitude compensation](#). In *AAAI Conference on Artificial Intelligence*.

Yung-Sung Chuang, Yujia Xie, Hongyin Luo, Yoon Kim, James Glass, and Pengcheng He. 2024. DoLa: Decoding by contrasting layers improves factuality in large language models. In *International Conference on Learning Representations (ICLR)*.

Christopher Clark, Kenton Lee, Ming-Wei Chang, Tom Kwiatkowski, Michael Collins, and Kristina Toutanova. 2019. BoolQ: Exploring the surprising difficulty of natural yes/no questions. In *Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.

Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. 2018. Think you have solved question answering? try ARC, the AI2 reasoning challenge. *arXiv preprint arXiv:1803.05457*.

Damai Dai, Li Dong, Yaru Hao, Zhifang Sui, Baobao Chang, and Furu Wei. 2022. Knowledge neurons in pretrained transformers. In *Annual Meeting of the Association for Computational Linguistics (ACL)*.

Hexuan Deng, Wenxiang Jiao, Xuebo Liu, Jing Li, Min Zhang, and Zhaopeng Tu. 2025. DRPruning: Efficient large language model pruning through distributionally robust optimization. In *Annual Meeting of the Association for Computational Linguistics (ACL)*.

Mostafa Elhoushi, Akshat Shrivastava, Diana Liskovich, Basil Hosmer, Bram Wasti, Liangzhen Lai, Anas Mahmoud, Bilge Acun, Saurabh Agarwal, Ahmed Roman, Ahmed A. Aly, Beidi Chen, and Carole-Jean Wu. 2024. [LayerSkip: Enabling early exit inference and self-speculative decoding](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*.

Mor Geva, Roei Schuster, Jonathan Berant, and Omer Levy. 2021. Transformer feed-forward layers are key-value memories. In <i>Conference on Empirical Methods in Natural Language Processing (EMNLP)</i> .	Alex Mallen, Akari Asai, Victor Zhong, Rajarshi Das, Daniel Khashabi, and Hannaneh Hajishirzi. 2023. When not to trust language models: Investigating effectiveness of parametric and non-parametric memories. In <i>Annual Meeting of the Association for Computational Linguistics (ACL)</i> .
Andrey Gromov, Kushal Tirumala, Hassan Shapourian, Paolo Glorioso, and Daniel A. Roberts. 2025. The unreasonable ineffectiveness of the deeper layers. In <i>International Conference on Learning Representations</i> .	Pere Martra. 2025. Fragile knowledge, robust instruction-following: The width pruning dichotomy in Llama-3.2. <i>arXiv preprint arXiv:2512.22671</i> .
Vipul Gupta, David Pantoja, Candace Ross, Adina Williams, and Megan Ung. 2024. Changing answer order can decrease MMLU accuracy. <i>arXiv preprint arXiv:2406.19470</i> .	Xin Men, Mingyu Xu, Qingyu Zhang, Bingning Wang, Hongyu Lin, Yaojie Lu, Xianpei Han, and Weipeng Chen. 2024. ShortGPT: Layers in large language models are more redundant than you expect. <i>arXiv preprint arXiv:2403.03853</i> .
Bowei He, Lihao Yin, Hui-Ling Zhen, Xiaokun Zhang, Mingxuan Yuan, and Chen Ma. 2025. PASER: Post-training data selection for efficient pruned large language model recovery. <i>arXiv preprint arXiv:2502.12594</i> .	Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. 2022. Locating and editing factual associations in GPT. In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> .
Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. 2021. Measuring massive multitask language understanding. In <i>International Conference on Learning Representations (ICLR)</i> .	Todor Mihaylov, Peter Clark, Tushar Khot, and Ashish Sabharwal. 2018. Can a suit of armor conduct electricity? a new dataset for open book question answering. In <i>Conference on Empirical Methods in Natural Language Processing (EMNLP)</i> .
Ajay Jaiswal, Zhe Gan, Xianzhi Du, Bowen Zhang, Zhangyang Wang, and Yinfei Yang. 2024. Compressing LLMs: The truth is rarely pure and never simple. In <i>International Conference on Learning Representations</i> .	Saurav Muralidharan, Sharath Turuvekere Sreenivas, Raviraj Joshi, Marcin Chochowski, Mostofa Patwary, Mohammad Shoenybi, Bryan Catanzaro, Jan Kautz, and Pavlo Molchanov. 2024. Compact language models via pruning and knowledge distillation. In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> .
Mandar Joshi, Eunsol Choi, Daniel Weld, and Luke Zettlemoyer. 2017. TriviaQA: A large scale distantly supervised challenge dataset for reading comprehension. In <i>Annual Meeting of the Association for Computational Linguistics (ACL)</i> .	Satya Sai Srinath Namburi, Makesh Sreedhar, Srinath Srinivasan, and Frederic Sala. 2023. The cost of compression: Investigating the impact of compression on parametric knowledge in language models. In <i>Findings of the Association for Computational Linguistics: EMNLP 2023</i> .
Bo-Kyeong Kim, Geonmin Kim, Tae-Ho Kim, Thibault Castells, Shinkook Choi, Junho Shin, and Hyoung-Kyu Song. 2024. Shortened LLaMA: Depth pruning for large language models with comparison of retraining methods. <i>arXiv preprint arXiv:2402.02834</i> .	nostalgebraist. 2020. Interpreting GPT: The logit lens. LessWrong. https://www.lesswrong.com/posts/AcKRB8wDpdaN6v6ru/interpreting-gpt-the-logit-lens .
Minkyu Kim, Vincent-Daniel Yun, Youngrae Kim, Suin Cho, Woosang Lim, and Sunwoo Lee. 2026. Re-thinking layer redundancy: Calibration matters more than search in LLM depth pruning. <i>arXiv preprint arXiv:2604.24938</i> .	OLMo Team, Pete Walsh, Luca Soldaini, Dirk Groeneveld, et al. 2025. 2 OLMo 2 Furious. <i>arXiv preprint arXiv:2501.00656</i> .
Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, Kristina Toutanova, Llion Jones, Matthew Kelcey, Ming-Wei Chang, Andrew M. Dai, Jakob Uszkoreit, Quoc Le, and Slav Petrov. 2019. Natural questions: A benchmark for question answering research. <i>Transactions of the Association for Computational Linguistics</i> , 7:453–466.	Denis Paperno, Germán Kruszewski, Angeliki Lazaridou, Quan Ngoc Pham, Raffaella Bernardi, Sandro Pezzelle, Marco Baroni, Gemma Boleda, and Raquel Fernández. 2016. The LAMBADA dataset: Word prediction requiring a broad discourse context. In <i>Annual Meeting of the Association for Computational Linguistics (ACL)</i> .
Yao Lu, Hao Cheng, Yujie Fang, Zeyu Wang, Jiaheng Wei, Dongwei Xu, Qi Xuan, Xiaoniu Yang, and Zhaowei Zhu. 2024. Reassessing layer pruning in LLMs: New insights and methods. <i>arXiv preprint arXiv:2411.15558</i> .	Keisuke Sakaguchi, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. 2021. WinoGrande: An adversarial Winograd schema challenge at scale. <i>Communications of the ACM</i> .

816	Maarten Sap, Hannah Rashkin, Derek Chen, Ronan	Mengzhou Xia, Tianyu Gao, Zhiyuan Zeng, and Danqi	871
817	Le Bras, and Yejin Choi. 2019. SocialIQA: Com-	Chen. 2024. Sheared LLaMA: Accelerating lan-	872
818	monsense reasoning about social interactions. In	guage model pre-training via structured pruning. In	873
819	<i>Conference on Empirical Methods in Natural Lan-</i>	<i>International Conference on Learning Representa-</i>	874
820	<i>guage Processing (EMNLP).</i>	<i>tions (ICLR).</i>	875
821	Boyuan Shi, Chang Liu, Chuanbao Gao, Xu Yang, and	Zhichao Xu, Ashim Gupta, Tao Li, Oliver Bentham, and	876
822	Xin Geng. 2026. Understanding performance col-	Vivek Srikumar. 2024. Beyond perplexity: Multi-	877
823	lapse in layer-pruned large language models via	dimensional safety evaluation of LLM compression.	878
824	decision representation transitions. <i>arXiv preprint</i>	In <i>Findings of the Association for Computational</i>	879
825	<i>arXiv:2605.07271.</i>	<i>Linguistics: EMNLP 2024.</i>	880
826	Shoaib Ahmed Siddiqui, Xin Dong, Greg Heinrich,	An Yang et al. 2025. Qwen3 technical report. <i>arXiv</i>	881
827	Thomas Breuel, Jan Kautz, David Krueger, and Pavlo	<i>preprint arXiv:2505.09388.</i>	882
828	Molchanov. 2024. A deeper look at depth pruning of	Yifei Yang, Zouying Cao, and Hai Zhao. 2024. LaCo:	883
829	LLMs. <i>arXiv preprint arXiv:2407.16286.</i>	Large language model pruning via layer collapse.	884
830	Luca Soldaini, Rodney Kinney, Akshita Bhagia, et al.	In <i>Findings of the Association for Computational</i>	885
831	2024. Dolma: An open corpus of three trillion to-	<i>Linguistics: EMNLP.</i>	886
832	kens for language model pretraining research. <i>arXiv</i>	Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali	887
833	<i>preprint arXiv:2402.00159.</i>	Farhadi, and Yejin Choi. 2019. HellaSwag: Can	888
834	Jiwon Song, Kyungseok Oh, Taesu Kim, Hyungjun	a machine really finish your sentence? In <i>Annual</i>	889
835	Kim, Yulhwa Kim, and Jae-Joon Kim. 2024. SLEB:	<i>Meeting of the Association for Computational Lin-</i>	890
836	Streamlining LLMs through redundancy verification	<i>guistics (ACL).</i>	891
837	and elimination of transformer blocks. In <i>Interna-</i>	Qingyu Zhang, Qianhao Yuan, Hongyu Lin, Yaojie Lu,	892
838	<i>tional Conference on Machine Learning (ICML).</i>	Xianpei Han, Le Sun, Xiang Li, Ming Xu, Jiarui	893
839	Sharath Turuvekere Sreenivas, Saurav Muralidharan,	Li, and Xiuyin Zhao. 2026. ShortOPD: Recovering	894
840	Raviraj Joshi, Marcin Chochowski, Ameya Sunil	pruned LLMs with short-to-long on-policy distilla-	895
841	Mahabaleshwarkar, Gerald Shen, Jiaqi Zeng, Zi-	tion. <i>arXiv preprint arXiv:2607.13124.</i>	896
842	jia Chen, Yoshi Suhara, Shizhe Diao, Chenhan Yu	Longguang Zhong, Fanqi Wan, Ruijun Chen, Xiaojun	897
843	Chen, Wei-Chun Chen, Hayley Ross, Oluwatobi	Quan, and Liangzhi Li. 2025. BlockPruner: Fine-	898
844	Olabiya, Ashwath Aithal, Oleksii Kuchaiev, Daniel	grained pruning for large language models. In <i>Find-</i>	899
845	Korzekwa, Pavlo Molchanov, Mostofa Patwary, and	<i>ings of the Association for Computational Linguistics:</i>	900
846	3 others. 2024. LLM pruning and distillation in	<i>ACL 2025</i> , pages 5065–5080.	901
847	practice: The minitron approach. <i>arXiv preprint</i>		
848	<i>arXiv:2408.11796.</i>		
849	Alon Talmor, Jonathan Herzig, Nicholas Lourie, and		
850	Jonathan Berant. 2019. CommonsenseQA: A ques-		
851	tion answering challenge targeting commonsense		
852	knowledge. In <i>Conference of the North American</i>		
853	<i>Chapter of the Association for Computational Lin-</i>		
854	<i>guistics (NAACL).</i>		
855	Shengkun Tang, Zekun Wang, Bo Zheng, Liangyu		
856	Wang, Rui Men, Siqi Zhang, Xiulong Yuan, Zi-		
857	han Qiu, Zhiqiang Shen, and Dayiheng Liu. 2026.		
858	SlimQwen: Exploring the pruning and distillation		
859	in large MoE model pre-training. <i>arXiv preprint</i>		
860	<i>arXiv:2605.08738.</i>		
861	Xinpeng Wang, Bolei Ma, Chengzhi Hu, Leon Weber-		
862	Genzel, Paul Röttger, Frauke Kreuter, Dirk Hovy, and		
863	Barbara Plank. 2024. “My Answer is C”: First-token		
864	probabilities do not match text answers in instruction-		
865	tuned language models. In <i>Findings of the Associa-</i>		
866	<i>tion for Computational Linguistics: ACL 2024.</i>		
867	Haryo Akbarianto Wibowo, Haiyue Song, Hideki		
868	Tanaka, Masao Utiyama, Alham Fikri Aji, and Raj		
869	Dabre. 2025. IteRABRe: Iterative recovery-aided		
870	block reduction. <i>arXiv preprint arXiv:2503.06291.</i>		

A Complete Results

The main paper emphasizes an observed-budget, single-run recovery gap. This appendix reports all consolidated checkpoints, including the qualitative 1B comparison, within-arm trajectories, and chance-adjusted summaries.

A.1 Held-out perplexity checkpoints

Table 4 distinguishes progressing checkpoints from true endpoints. The fully random-init arm is a random language model with the same 16-layer architecture: its decoder blocks, embedding, final norm, and output head are initialized without OLMo-2 weights. It uses a different learning rate, so it is a budgeted operating point rather than an initialization-only ablation. Table 5 reports both immediate damage and the healed non-contiguous endpoint. Its 16 retained layers are selected by cosine block influence on 128 Dolmino windows and include the original final layer.

A.2 Within-architecture training trajectories

The complete 11-task keep8 trajectory in Table 6 gives a within-arm view over available checkpoints. From 44k to 121k, most reported tasks improve, while aggregate MMLU changes from .2463 to .2535 and remains near chance. Because no common per-item 44k prediction set or replicated training run is available, we use this only as descriptive evidence against a uniform late gain over the observed interval, not against recovery after more training.

The OLMo-2-0425-1B arm shows the same qualitative separation (Figure 5; Tables 7–8): PPL falls 17.619 $\rightarrow$ 15.628 and HellaSwag/ARC-Easy/PIQA improve, while MMLU is .2495/.2558/.2529/.2480. Because this is one more-compressed arm from the same model family with a weaker base, we treat it as qualitative context, not an independent replication.

A.3 Capability recovery and continued healing

Table 11 exposes why a single average benchmark score would be misleading. Fully random initialization nearly matches inherited keep14 on ARC-Easy, PIQA, and BoolQ, but inherited initialization retains more signal on WinoGrande, MMLU, and several other downstream tasks. The late-healing comparison in Table 13 shows no broad rollback through 200k: from 153.5k to 200k, PPL improves

Scale	Arm	step	inherited	total layers	PPL↓	tax↓
7B	full base	—	32	32	7.398	1.000×
	keep8+fresh2	5k	8	10	22.331	3.019×
	keep8+fresh2	15k	8	10	17.868	2.416×
	keep8+fresh2	25k	8	10	16.426	2.220×
	keep8+fresh2	35k	8	10	15.612	2.110×
	keep8+fresh2	44k	8	10	15.131	2.045×
	keep8+fresh2	121k [†]	8	10	13.333	1.802×
	keep10+fresh2	10k	10	12	17.239	2.330×
	keep10+fresh2	83.5k [†]	10	12	12.816	1.732×
	keep12+fresh2	111.5k	12	14	11.566	1.563×
	keep12+fresh2	124k [†]	12	14	11.443	1.547×
	keep14+fresh2	128k	14	16	10.826	1.463×
	keep14+fresh2	153.5k	14	16	10.693	1.445×
	keep14+fresh2	200k	14	16	10.561	1.428×
	full32 continued@25k	25k [‡]	32	32	7.670	1.037×
	ShortGPT-16	200k	16	16	9.780	1.322×
	frozen-front control	200k	14	16	12.797	1.730×
	fully random-init control	200k	0	16	11.498	1.554×
1B	full base	—	16	16	10.642	1.000×
	keep7+fresh2	50k	7	9	17.619	1.656×
	keep7+fresh2	100k	7	9	16.161	1.519×
	keep7+fresh2	147k	7	9	15.628	1.469×

Table 4: **All held-out perplexity checkpoints used in this study.** Every point uses the same disjoint 4096 $\times$ 2048 Dolmino validation shard (8,384,512 target tokens) and token-weighted eight-shard merging. The frozen-front arm retains the same pretrained prefix but does not update it; the fully random-init arm has the same 16-layer shape but loads no pretrained weights. ShortGPT inherits non-contiguous layers [0–12, 16, 17, 31] and appends no fresh blocks. [†]The available endpoint values come from the retained keep8 at 121k, keep10 at 83.5k, and keep12 at 124k checkpoints after metric-based stopping; no literal step-200k evaluation is available for these arms, and PPL was still decreasing. [‡]All reported full32 results use its available step-25k checkpoint; it is not a long-horizon matched counterfactual.

by 0.132 while MMLU gains only 0.67 points and most other metrics move by less than one point. Table 14 additionally exposes raw and continuation-length-normalized scoring for BoolQ, CommonsenseQA, and SocialIQA, whose options have unequal lengths; the main tables consistently use raw accuracy for these tasks. Tables 16 and 17 report the complete MMLU interface control and the zero-shot closed-book results used in the main analysis. Table 15 uses a common per-item MMLU rerun for the three same-shape 200k operating points. All pairwise intervals exclude zero, but these are evaluation-item intervals rather than training-seed uncertainty and the rerun aggregates differ slightly from separately reported headline cells.

The confidence intervals in Table 12 distinguish the available checkpoints conditionally: keep8 and fully random init include chance, while keep10, keep12, and keep14 are above it. These are not between-training-run intervals. The subject-group table below describes heterogeneity rather than testing domain-specific effects.

ShortGPT point	PPL	core6	MMLU
step 0	401.124	—	.2620
step 200k	9.780	.6215	.4739

Table 5: **Non-contiguous ShortGPT endpoint.** Cosine block influence selects 16 of 32 pretrained layers, [0–12, 16, 17, 31], with no fresh tail. The unhealed model is initially degenerate; after the same 200k-step Dolmino healing budget it recovers 63.0% of base above-chance MMLU.

Task	metric	10k	25k	44k	121k	$\Delta_{44 \rightarrow 121}$
MMLU	raw	.2542	.2502	.2463	.2535	+0.72
HellaSwag	norm.	.3915	.4390	.4694	.5169	+4.75
ARC-Challenge	norm.	.2654	.3114	.3140	.3601	+4.61
ARC-Easy	norm.	.5581	.5930	.6103	.6519	+4.16
PIQA	norm.	.6632	.6801	.6937	.7165	+2.28
WinoGrande	raw	.5209	.5083	.5185	.5312	+1.27
OpenBookQA	norm.	.3000	.3360	.3320	.3660	+3.40
LAMBADA	raw	.3429	.3827	.4333	.4463	+1.31
BoolQ	raw	.5713	.6080	.5881	.5948	+0.67
CommonsenseQA	raw	.3923	.4103	.4423	.4521	+0.98
SocialIQA	raw	.3767	.3874	.4002	.3976	−0.26

Table 6: **Complete 11-task keep8 downstream trajectory.** The final column is the accuracy-point change from 44k to the retained 121k endpoint checkpoint under the reported metric shown in column two; “norm.” denotes accuracy based on token-normalized candidate log-likelihood. MMLU remains near chance while most other tasks improve. Earlier PPL checkpoints use a separate grid; the final held-out PPL is reported at 121k in Table 4. This table reports aggregate checkpoints, not a paired 44k–121k test or a replicated trajectory.

A.4 MMLU subject-group results

Table 18 quantifies the broad-domain pattern. The inherited keep14 model is strongest in social science and the heterogeneous “Other” group, whereas science, technology, engineering, and mathematics (STEM) and humanities remain closer to chance. Table 22, placed at the end of the appendix as two side-by-side panels, reports every one of the 57 subjects at the 153.5k and final 200k keep14 checkpoints.

B Evaluation and Reproducibility Details

B.1 Data, optimization, and architecture reconstruction

All training uses English-language OLMo-2/Dolmino artifacts. The OLMo-2 checkpoint and Dolmino release are used under their published research licenses; the evaluation datasets retain their upstream terms and are not redistributed with our artifacts. Experiments ran on 8-GPU H20 or B200 nodes depending on the arm. The

Model	step	HS	ARC-C	ARC-E	PIQA	WinoG	OBQA	PPL
1B full base	—	.683	.422	.735	.757	.643	.400	10.642
keep7+fresh2	50k	.411	.289	.568	.671	.531	.318	17.619
keep7+fresh2	100k	.438	.300	.593	.687	.526	.326	16.161
keep7+fresh2	147k	.450	.311	.595	.692	.530	.326	15.628
keep7+fresh2	148.5k	.448	.294	.606	.695	.544	.330	—

Table 7: **OLMo-2-1B core-task trajectory.** We report `acc_norm` (continuation-length-normalized likelihood accuracy) except for WinoGrande, where raw and normalized accuracy coincide. Abbreviations: HS = HellaSwag; ARC-C/E = ARC-Challenge/Easy; WinoG = WinoGrande; OBQA = OpenBookQA; PPL = perplexity.

Model	step	MMLU	LAMBADA	BoolQ	CSQA	SIQA
1B full base	—	.3820	.6480	.624	.553	.422
keep7+fresh2	50k	.2495	.3536	.618	.368	.384
keep7+fresh2	100k	.2558	.3740	.597	.380	.394
keep7+fresh2	147k	.2529	.4021	.584	.391	.384
keep7+fresh2	150k	.2480	.3969	.577	.387	.386

Table 8: **OLMo-2-1B knowledge/comprehension trajectory.** MMLU remains at the .25 chance floor across the complete measured trajectory, whereas LAMBADA and commonsense tasks improve. Raw accuracy is shown for these five tasks.

principal keep14 run used 200k optimizer steps; the full32 branch has an available 25k checkpoint, and shallower prefix arms use the literal latest retained checkpoints stated in the tables. These were not selected by a registered common stopping rule. Because exploratory and failed runs were not logged uniformly, we cannot provide an exact project-wide GPU-hour total. We therefore report per-run hardware and steps rather than an unreliable aggregate.

Healing uses the DataComp for Language Models (DCLM) portion of `dolmino-mix-1124`, tokenized into 2048-token windows with the OLMo-2 tokenizer. The training array is disjoint from the fixed validation shard of 4096 windows. All main arms use effective batch size 128 and a cosine schedule after 150 warmup steps. In the executed distributed runs, all parameters in inherited arms use peak/minimum learning rates $2 \times 10^{-5} / 2 \times 10^{-6}$; the fully random-init arm uses $10^{-4} / 10^{-5}$ for all parameters. We use AdamW with $\beta_1=0.9$, $\beta_2=0.95$, $\epsilon=10^{-8}$, weight decay 0.1 on matrix parameters and zero on vectors, and global gradient-norm clipping at 1.0. Parameters and optimizer states use 32-bit floating point (fp32), while forward passes use bfloat16 (bf16) automatic mixed precision; gradient checkpointing is enabled. The maximum is 200k optimizer steps. No BA-




Figure 5: **Qualitative OLMo-2-1B trajectory.** PPL and several ordinary tasks improve while MMLU stays near chance over the observed keep7 interval.

Qwen3-8B arm	Depth	PPL↓	MMLU letter
Full base	36/36	11.42	.7297
keep12 + fresh2, 200k	14/36	23.49	.2495

Table 9: **Cross-family endpoint diagnostic.** The Qwen3-8B (Yang et al., 2025) arm heals on SlimPajama rather than Dolmino and is more compressed than the principal OLMo arm. It supports only the directional answer-letter finding; absolute PPL and recovery magnitude are not compared across families. This endpoint does not include the content-MMLU and closed-book diagnostics used in the principal study.

BILong or other synthetic task data is mixed into healing.

Runs do not set an explicit random seed, so the exact fresh initialization and initial shuffle are not reproducible from the command line alone. The keep14 run resumed once from step 34.5k with model, optimizer, and pseudorandom-number generator states restored. The checkpoint did not store the within-epoch data-loader offset, so the resumed iterator restarted that epoch’s distributed shuffle; token counts in this paper are therefore nominal presentations rather than claims about unique examples.

For each inherited arm, launch-time checks require exact equality between the base and transplanted front blocks, embedding, final norm, and language-model (LM) vocabulary output head ($\max |\theta_{\text{arm}} - \theta_{\text{base}}| = 0$). Fresh-block norms are initialized to one and fresh query projections have standard deviation 0.02. Evaluation reconstructs the pruned architecture from checkpoint metadata

and the saved parameter-state dictionary, then loads every tensor strictly. The keep14 models have 4.060B parameters.

B.2 Likelihood scoring and dataset sizes

For PopQA, TriviaQA, and NQ-open, evaluation is zero-shot and uses no retrieval. The prompt is `Question: {q}\nAnswer:;` generation is greedy with no chat template or inserted BOS and a maximum of 32 new tokens. The prediction is the first completion line. Answers are lowercased, stripped of punctuation/articles/extra whitespace, and scored against all available aliases. PopQA’s headline is normalized-answer containment; TriviaQA and NQ-open use exact match. We additionally retain exact match, containment, and token F1 for sensitivity analysis.

For each multiple-choice candidate, the evaluator concatenates context and continuation, supplies the candidate tokens to the model, and sums their conditional log probabilities. This teacher-forced scoring never uses free generation. For the broad likelihood suite, `acc_norm` follows the task harness’s continuation-length convention. The separate MMLU content protocol reports both summed log-likelihood and mean log-likelihood per continuation token; the latter is the content headline. Winogrande uses its standard cloze format, in which each candidate fills the blank in the supplied sentence. MMLU candidates are answer letters under a fixed subject-description prompt. LAMBADA reports whether the entire final-word continuation is the greedy token sequence.

Evaluation integrity. Perplexity is merged as $\exp(\sum_g \text{NLL}_g / \sum_g n_g)$ over eight shards, where NLL_g is shard g ’s summed negative log-likelihood and n_g is its token count; we never average per-shard perplexities. At keep14 steps 128k, 153.5k, and 200k, independent recomputation from all shard files reproduces every PPL and downstream cell to 10^{-9} . All evaluated task cells have zero NaNs and full sample counts. BoolQ contains the same two overlength passages in every run containing BoolQ; they are truncated consistently and still scored. No other task is truncated. The full 7B base reproduces expected OLMo-2 reference values (HellaSwag .805, MMLU .6053, LAMBADA .732), and the 1B base similarly reproduces its expected range.

Arm	step	MMLU	HS	ARC-C	ARC-E	PIQA	WinoG	OBQA	LAMB.	BoolQ	CSQA	SIQA
7B base full (32L)	—	.6053	.805	.571	.829	.811	.744	.462	.732	.815	.665	.502
keep8 (10L)	121k	.2535	.517	.360	.652	.717	.531	.366	.446	.595	.452	.398
keep10 (12L)	83.5k	.2718	.547	.366	.644	.726	.542	.356	.496	.609	.454	.415
keep12 (14L)	124k	.2752	.599	.394	.689	.737	.607	.376	.535	.587	.496	.410
keep14 (16L)	128k	.3012	.631	.426	.702	.747	.630	.402	.575	.639	.505	.423
keep14 (16L)	153.5k	.3124	.643	.442	.705	.745	.633	.406	.570	.606	.506	.441
keep14 (16L)	200k	.3191	.645	.438	.705	.745	.626	.404	.577	.638	.499	.434
ShortGPT-16	200k	.4739	.685	.476	.746	.758	.655	.408	.619	.729	.534	.442
frozen-front (16L)	200k	.2628	.595	.381	.666	.724	.646	.366	.513	.592	.453	.414
fully random-init (16L)	200k	.2461	.578	.414	.697	.733	.545	.384	.484	.614	.450	.416

Table 10: **Downstream zero-shot likelihood scoring (7B)**. `acc_norm` is used for HS, ARC-C/E, PIQA, and OBQA; raw accuracy is used for MMLU, WinoG, LAMBADA, BoolQ, CSQA, and SIQA. Here `acc_norm` means accuracy from continuation-length-normalized log-likelihood. Appendix Table 14 reports both raw and normalized views for the final three tasks. Abbreviations: HS = HellaSwag; ARC-C/E = ARC-Challenge/Easy; WinoG = WinoGrande; OBQA = OpenBookQA; LAMB. = LAMBADA; CSQA = CommonsenseQA; SIQA = SocialQA. The prefix ladder places keep8 at MMLU chance and final keep14 at 19.4% above-chance recovery. ShortGPT-16 reaches 63.0%, but also inherits two more pretrained blocks and the original final layer. Frozen-front recovers only 3.6%, while fully random init remains at chance despite nontrivial scores elsewhere and similar raw BoolQ accuracy to keep14 (.614 vs. .638). The shallow prefix ladder uses unequal latest retained checkpoints and is not step- or compute-matched; keep14, ShortGPT, and the same-shape operating points are literal 200k evaluations.

C Supplementary Layer-wise Readout Background

This section is not part of the evidence for the recovery-path claims. It documents exploratory readouts on intact models to provide background only; the thresholds use different datasets, probes, and definitions and are neither replicated mechanism estimates nor causal storage localizations. The analysis follows logit- and tuned-lens work (nostalgebraist, 2020; Belrose et al., 2023), alongside descriptive work on feed-forward memory, knowledge neurons, causal tracing, contrastive decoding, and early exit (Geva et al., 2021; Dai et al., 2022; Meng et al., 2022; Chuang et al., 2024; Elhoushi et al., 2024).

The MMLU logit lens uses 1,000 MMLU test questions. At each layerwise hidden-state depth, the final-prompt-position hidden state is passed through the model’s final normalization and output head; accuracy is the argmax among the letter tokens {A,B,C,D}. We define onset as the first layer exceeding chance by .05 and sat_{95} as the first layer reaching 95% of the top-layer above-chance signal. On OLMo these occur at layers 18 and 19, while sat_{99} is reached at layer 27. Complete OLMo and supplemental Qwen values appear in Table 21.

An exploratory same-model OLMo probe adds frozen linear classifiers for WiC, SST-2, and RTE at every depth plus exact next-token top-1 agreement with the final layer. Their final-layer-relative sat_{95} depths are layers 1, 5, and 1, giving a mean of $2.33/32 = 0.073L$; because intermediate probes can outperform the final layer, this is a readout




Figure 6: **Probe-dependent readout depths on OLMo-2-7B**. The shaded band marks prefix cuts used by the available behavioral points. Thresholds use different readouts and are descriptive; they are neither directly comparable causal boundaries nor locations of stored knowledge.

heuristic rather than a causal semantic boundary. Next-token agreement reaches sat_{95} only at layer 32 (1.000L). These descriptive thresholds form the ordering semantic < MMLU < next-token under the stated definitions, but they are not used to explain keep14 or ShortGPT.

Family	Task	chance	base	inherited	random init	recovery (%)
				reported score		inherited / random
surface	ARC-Easy	.25	.829	.705	.697	78.6/77.2
surface	PIQA	.50	.811	.745	.733	78.9/74.9
reasoning	HellaSwag	.25	.805	.645	.578	71.1/59.1
reasoning	ARC-Challenge	.25	.571	.438	.414	58.5/51.1
reasoning	WinoGrande	.50	.744	.626	.545	51.6/18.4
reasoning	OpenBookQA	.25	.462	.404	.384	72.6/63.2
reasoning	LAMBADA	.00	.732	.577	.484	78.9/66.1
reasoning	SocialQA	.333	.502	.434	.416	59.8/49.1
reasoning	CommonsenseQA	.20	.665	.499	.450	64.3/53.9
in-context	BoolQ	.50	.815	.638	.614	43.9/36.2
answer-letter	MMLU	.25	.6053	.3191	.2461	19.4/ - 1.1

Table 11: **Exact above-chance recovery for inherited keep14 and the fully random-init control at the matched step 200k.** Recovery is $100(x - c)/(b - c)$ for score x , chance c , and full-base score b . Family labels are descriptive: “surface” contains ARC-Easy and PIQA; “in-context” denotes passage-grounded BoolQ; “answer-letter” denotes the MMLU protocol; the remaining tasks are grouped descriptively as “reasoning.” Reported scores use continuation-length-normalized likelihood for HellaSwag, ARC, PIQA, and OpenBookQA and raw accuracy otherwise. The negative random-init MMLU value means statistically chance-level performance. The random-init arm also randomizes lexical modules and uses a different learning rate, so the table is an operating-point comparison rather than an initialization-only or decoder-block-only ablation. Negative recovery denotes a score below the chance-adjusted reference.

Arm	MMLU	SE	approx. 95% CI
full 7B base	.6053	.0041	[.5972, .6134]
keep8, 121k	.2535	.0037	[.2463, .2607]
keep10, 83.5k	.2718	.0038	[.2644, .2792]
keep12, 124k	.2752	.0038	[.2678, .2826]
keep14, 128k	.3012	.0039	[.2936, .3088]
keep14, 153.5k	.3124	.0039	[.3047, .3201]
keep14, 200k	.3191	.0039	[.3114, .3268]
ShortGPT-16, 200k	.4739	.0042	[.4656, .4822]
frozen-front, 200k	.2628	.0037	[.2555, .2701]
fully random init, 200k	.2461	.0036	[.2390, .2532]

Table 12: **Marginal MMLU uncertainty across the 7B frontier.** Standard errors and Wald intervals use $n = 14,042$ questions. keep8 and fully random init include the .25 chance floor; keep10/keep12, frozen-front, all keep14 checkpoints, and ShortGPT are above it. These marginal intervals summarize each arm separately; Table 15 gives the stronger item-paired three-control comparisons.

Metric	128k	153.5k	200k
held-out PPL	10.826	10.693	10.561
MMLU	.3012	.3124	.3191
HellaSwag	.631	.643	.645
ARC-Challenge	.426	.442	.438
ARC-Easy	.702	.705	.705
PIQA	.747	.745	.745
WinoGrande	.630	.633	.626
OpenBookQA	.402	.406	.404
LAMBADA	.575	.570	.577
BoolQ (raw acc.)	.639	.606	.638
CommonsenseQA (raw)	.505	.506	.499
SocialQA (raw)	.423	.441	.434

Table 13: **Complete late keep14 trajectory.** From 153.5k to 200k, held-out PPL improves by 0.132 while MMLU gains only 0.67 points. Over the paired 14,042-item rerun from 128k to 200k, MMLU increases by 1.68 points (95% CI [1.08, 2.29], exact McNemar $p = 4.12 \times 10^{-8}$). Most other metrics remain within one point from 153.5k to 200k, with mixed signs; continued healing yields a small but detectable answer-letter gain rather than broad late catch-up.

Task	metric	base	k8	k12	k14-p	k14-f	frozen	random
BoolQ	raw	.815	.588	.610	.606	.638	.592	.614
BoolQ	norm.	.753	.620	.657	.682	.689	.700	.619
CSQA	raw	.665	.442	.508	.506	.499	.453	.450
CSQA	norm.	.652	.405	.471	.479	.476	.467	.405
SIQA	raw	.502	.400	.415	.441	.434	.414	.416
SIQA	norm.	.547	.431	.460	.474	.474	.451	.443

Table 14: **Raw versus continuation-length-normalized accuracy.** The normalized score divides candidate log-likelihood by continuation character length before selecting an answer. The main results report raw accuracy for BoolQ, CommonsenseQA, and SocialIQA. Normalization can materially change BoolQ and SIQA, so both views are reported here. k14-p/k14-f denote keep14 at 153.5k/200k; frozen-front and random init are the other 200k operating points.

Comparison	n	Δ (pp)	McNemar p	95% CI (pp)
keep14 – random init	14,042	+7.11	1.63×10^{-46}	[6.14, 8.09]
keep14 – frozen front	14,042	+5.50	6.99×10^{-27}	[4.51, 6.49]
frozen front – random init	14,042	+1.61	2.59×10^{-3}	[0.58, 2.64]

Table 15: **Paired MMLU comparisons among the three 200k, 16-layer operating points.** Δ follows the row order and is computed from the aligned per-item rerun (keep14 .31797, frozen .26293, random .24683), which differs slightly from the separately rounded aggregate cells. McNemar uses an exact two-sided binomial test on discordant correctness pairs; 95% confidence intervals use 10,000 paired bootstrap resamples with seed 1234. All intervals exclude zero.

Arm	letter	content raw	content norm.
Full base	.6054	.4400	.4706
Full32 at 25k	.5877	.4344	.4662
keep8 endpoint [†]	.2550	.3219	.3423
keep10 endpoint [†]	.2720	.3230	.3445
keep12 endpoint [†]	.2728	.3419	.3629
keep14 at 200k	.3184	.3548	.3832
Frozen front at 200k	.2624	.3349	.3604
Random init at 200k	.2470	.3439	.3598
ShortGPT-16 at 200k	.4742	.3764	.4012

Table 16: **MMLU scoring-interface control on all 14,042 items.** Letter scores one of A/B/C/D after displaying the options. Content scores the complete option text. “Content raw” selects by summed continuation log-likelihood; “content norm.” selects by mean log-likelihood per continuation token and is the reported content headline. The two protocols change prompt, candidate string, tokenization, and normalization together and have no paired uncertainty analysis. [†]The shallow rows use retained keep8 at 121k, keep10 at 83.5k, and keep12 at 124k; no literal step-200k content evaluations are available for these arms.

Arm	PopQA contains	TriviaQA EM	NQ EM
Full base	.2571	.6355	.2050
Full32 at 25k	.2280	.5715	.1582
keep14 at 200k	.1415	.2940	.0598
Frozen front at 200k	.1283	.2477	.0496
Random init at 200k	.1112	.2086	.0632
keep8 endpoint [†]	.1246	.1577	.0368

Table 17: **Zero-shot, no-retrieval closed-book evaluation.** All arms use the same plain Question:/Answer: generation protocol. All full32 cells use the intact 32-layer step-25k checkpoint. [†]keep8 has a different 10-layer shape; its retained checkpoint is step 121k and it is included only as a depth reference.

Group	n	base	full32	keep14	ShortGPT	random
STEM	3,864	.534	.507	.294	.405	.236
Humanities	4,705	.555	.542	.299	.428	.252
Social science	3,077	.709	.690	.335	.554	.240
Other	2,396	.686	.670	.377	.574	.259

Table 18: **Sample-weighted MMLU broad-group accuracy.** full32 is the 25k continued-full-model control; keep14, ShortGPT, and random are their reported 200k endpoints. Group definitions follow the standard four-way MMLU taxonomy.

Task	split / mode	n	chance	metric
HellaSwag	validation, 4-choice	10,042	.25	acc_norm
ARC-Challenge	test, 4-choice	1,172	.25	acc_norm
ARC-Easy	test, 4-choice	2,376	.25	acc_norm
PIQA	validation, 2-choice	1,838	.50	acc_norm
WinoGrande	validation, 2-choice cloze	1,267	.50	accuracy
OpenBookQA	test, 4-choice	500	.25	acc_norm
MMLU	test, 57 subjects, 4-choice letters	14,042	.25	accuracy
LAMBADA-OpenAI	test, greedy final word	5,153	≈ 0	exact greedy accuracy
BoolQ	validation, yes/no	3,270	.50	accuracy
CommonsenseQA	validation, 5-choice	1,221	.20	accuracy
SocialIQA	validation, 3-choice	1,954	.333	accuracy

Table 19: **Evaluation suite and chance floors.** All tasks are zero-shot and likelihood scored; no free-form generation or chat template is used. acc_norm selects answers after dividing candidate log-likelihood by continuation character length. All listed examples are evaluated.

OLMo-2-7B quantity	layer	fractional depth
semantic probe mean sat ₉₅	2.33	0.073 L
MMLU onset (chance+0.05)	18	0.562 L
MMLU sat ₉₅	19	0.594 L
next-token sat ₉₅	32	1.000 L

Table 20: **Three ordered readout depths on the same OLMo model.** Semantic sat₉₅ averages WiC, SST-2, and RTE depths measured relative to each final-layer probe score. MMLU uses the final norm and output head; next-token sat₉₅ is exact top-1 agreement with the final layer. These are readout thresholds, not causal boundaries.

(a) OLMo-2-7B								(b) Qwen3-8B							
L	frac.	acc.	gold LL	L	frac.	acc.	gold LL	L	frac.	acc.	gold LL	L	frac.	acc.	gold LL
0	0.0000	0.230	-14.499	17	0.5312	0.251	-11.947	0	0.0000	0.272	-22.017	19	0.5278	0.237	-13.084
1	0.0312	0.232	-14.024	18	0.5625	0.326	-11.652	1	0.0278	0.234	-24.014	20	0.5556	0.237	-12.790
2	0.0625	0.276	-13.244	19	0.5938	0.544	-8.517	2	0.0556	0.240	-24.741	21	0.5833	0.236	-12.816
3	0.0938	0.240	-12.939	20	0.6250	0.541	-8.038	3	0.0833	0.236	-24.112	22	0.6111	0.235	-10.357
4	0.1250	0.272	-13.112	21	0.6562	0.538	-7.002	4	0.1111	0.233	-22.029	23	0.6389	0.232	-10.675
5	0.1562	0.259	-12.793	22	0.6875	0.542	-5.694	5	0.1389	0.237	-17.364	24	0.6667	0.236	-9.027
6	0.1875	0.262	-12.206	23	0.7188	0.540	-6.038	6	0.1667	0.242	-12.939	25	0.6944	0.621	-6.590
7	0.2188	0.228	-11.715	24	0.7500	0.535	-4.847	7	0.1944	0.234	-12.999	26	0.7222	0.625	-6.928
8	0.2500	0.223	-11.116	25	0.7812	0.530	-5.091	8	0.2222	0.229	-13.867	27	0.7500	0.624	-6.054
9	0.2812	0.224	-11.294	26	0.8125	0.531	-5.701	9	0.2500	0.241	-14.885	28	0.7778	0.635	-6.225
10	0.3125	0.236	-11.493	27	0.8438	0.548	-5.345	10	0.2778	0.237	-14.034	29	0.8056	0.572	-6.685
11	0.3438	0.248	-12.051	28	0.8750	0.549	-5.456	11	0.3056	0.229	-14.021	30	0.8333	0.629	-7.052
12	0.3750	0.268	-11.847	29	0.9062	0.545	-5.402	12	0.3333	0.228	-13.218	31	0.8611	0.630	-7.164
13	0.4062	0.228	-11.499	30	0.9375	0.533	-5.338	13	0.3611	0.225	-13.516	32	0.8889	0.629	-7.869
14	0.4375	0.235	-11.295	31	0.9688	0.540	-4.475	14	0.3889	0.229	-13.597	33	0.9167	0.629	-7.898
15	0.4688	0.249	-11.555	32	1.0000	0.551	-3.791	15	0.4167	0.231	-13.184	34	0.9444	0.638	-7.334
16	0.5000	0.234	-11.960					16	0.4444	0.237	-14.466	35	0.9722	0.636	-4.900
								17	0.4722	0.236	-14.439	36	1.0000	0.632	-5.412
								18	0.5000	0.231	-13.220				

Table 21: **Complete MMLU logit-lens trajectories.** Layer 0 is the embedding/input residual state and layer L follows decoder block L . Accuracy is four-letter MMLU accuracy; gold LL is full-vocabulary gold-letter log-likelihood on the same 1,000 questions. These are output-head readouts, not causal storage localizations.

Subject	n	base	k8	k12	k14-p	k14-f	frozen	rand	Subject	n	base	k8	k12	k14-p	k14-f	frozen	rand
abstract_algebra	100	.310	.300	.240	.260	.270	.310	.240	high_school_statistics	216	.514	.171	.338	.204	.213	.218	.194
anatomy	135	.622	.289	.296	.363	.363	.296	.215	high_school_us_history	204	.824	.314	.230	.319	.304	.299	.309
astronomy	152	.730	.171	.289	.316	.322	.283	.178	high_school_world_history	237	.797	.253	.270	.325	.354	.274	.253
business_ethics	100	.590	.250	.340	.390	.430	.260	.220	human_aging	223	.650	.283	.139	.368	.359	.265	.341
clinical_knowledge	265	.649	.242	.275	.317	.317	.257	.249	human_sexuality	131	.756	.221	.298	.382	.382	.298	.321
college_biology	144	.771	.278	.271	.347	.326	.250	.264	international_law	121	.777	.248	.240	.281	.264	.355	.240
college_chemistry	100	.470	.190	.310	.270	.230	.260	.200	jurisprudence	108	.685	.204	.269	.287	.315	.231	.259
college_computer_science	100	.490	.240	.300	.300	.290	.240	.270	logical_fallacies	163	.693	.233	.215	.380	.350	.209	.221
college_mathematics	100	.360	.200	.310	.240	.320	.270	.230	machine_learning	112	.455	.170	.277	.375	.330	.277	.321
college_medicine	173	.566	.277	.341	.295	.289	.249	.214	management	103	.777	.204	.359	.320	.369	.282	.194
college_physics	102	.422	.196	.304	.225	.275	.265	.147	marketing	234	.825	.286	.274	.440	.440	.286	.312
computer_security	100	.680	.230	.250	.410	.400	.340	.290	medical_genetics	100	.740	.170	.300	.370	.400	.340	.320
conceptual_physics	235	.570	.238	.264	.311	.332	.243	.277	miscellaneous	783	.803	.271	.278	.436	.457	.268	.244
econometrics	114	.360	.246	.193	.246	.246	.325	.237	moral_disputes	346	.694	.225	.260	.329	.321	.272	.237
electrical_engineering	145	.600	.255	.276	.359	.366	.345	.248	moral_scenarios	895	.279	.238	.238	.244	.253	.244	.238
elementary_mathematics	378	.397	.278	.254	.296	.272	.259	.222	nutrition	306	.693	.239	.307	.324	.327	.320	.242
formal_logic	126	.389	.302	.325	.198	.262	.183	.302	philosophy	311	.711	.235	.270	.322	.318	.257	.222
global_facts	100	.360	.340	.330	.280	.290	.320	.190	prehistory	324	.741	.318	.275	.346	.358	.299	.238
high_school_biology	310	.765	.223	.319	.319	.319	.239	.190	professional_accounting	282	.415	.266	.245	.238	.241	.262	.280
high_school_chemistry	203	.493	.261	.246	.256	.232	.202	.167	professional_law	1534	.460	.244	.246	.259	.272	.245	.263
high_school_computer_science	100	.590	.290	.220	.300	.290	.280	.290	professional_medicine	272	.570	.199	.239	.228	.232	.301	.301
high_school_european_history	165	.752	.236	.236	.242	.309	.194	.206	professional_psychology	612	.611	.243	.247	.320	.302	.260	.275
high_school_geography	198	.783	.242	.308	.384	.328	.288	.207	public_relations	110	.627	.282	.236	.291	.327	.309	.236
high_school_government_and_politics	193	.870	.202	.285	.368	.358	.244	.197	security_studies	245	.727	.367	.339	.294	.335	.245	.245
high_school_macroeconomics	390	.636	.226	.338	.295	.321	.285	.221	sociology	201	.836	.289	.279	.418	.428	.194	.294
high_school_mathematics	270	.296	.230	.278	.281	.248	.259	.219	us_foreign_policy	100	.850	.230	.340	.480	.450	.270	.340
high_school_microeconomics	238	.655	.239	.269	.277	.277	.252	.185	virology	166	.524	.241	.313	.355	.386	.283	.313
high_school_physics	151	.404	.238	.358	.225	.245	.298	.212	world_religions	171	.836	.246	.292	.491	.509	.287	.322
high_school_psychology	545	.807	.196	.297	.317	.358	.229	.207									

Table 22: **Complete 57-subject MMLU results.** All values are accuracy. k14-p and k14-f are inherited keep14 at 153.5k and 200k; frozen-front and random init are the two other 200k same-shape operating points. Subjects are split alphabetically across the two panels.